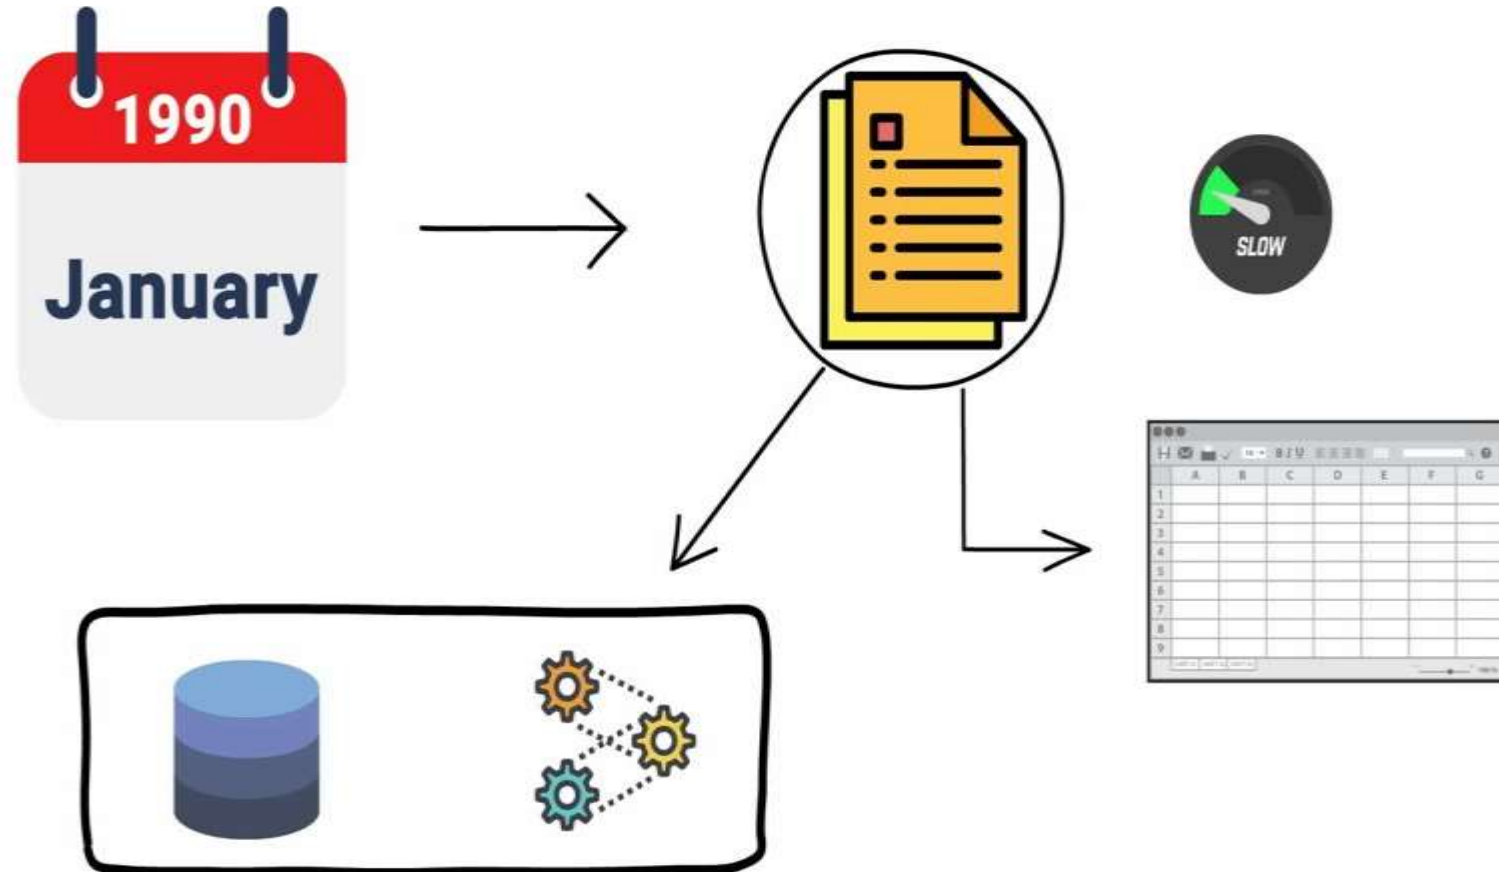


Hadoop


The rapid evolution of the internet has led to an explosion of diverse data types being generated at unprecedented speeds, surpassing the capabilities of traditional storage and processing methods.




The emergence of semi-structured and unstructured data, such as emails, images, audio, and video, necessitated the use of multiple storage units and processors to effectively manage big data.



The solution to insufficient storage and processing power was the implementation of multiple storage units and processors, leading to the development of Hadoop, which efficiently stores and processes large amounts of data using a cluster of commodity hardware.

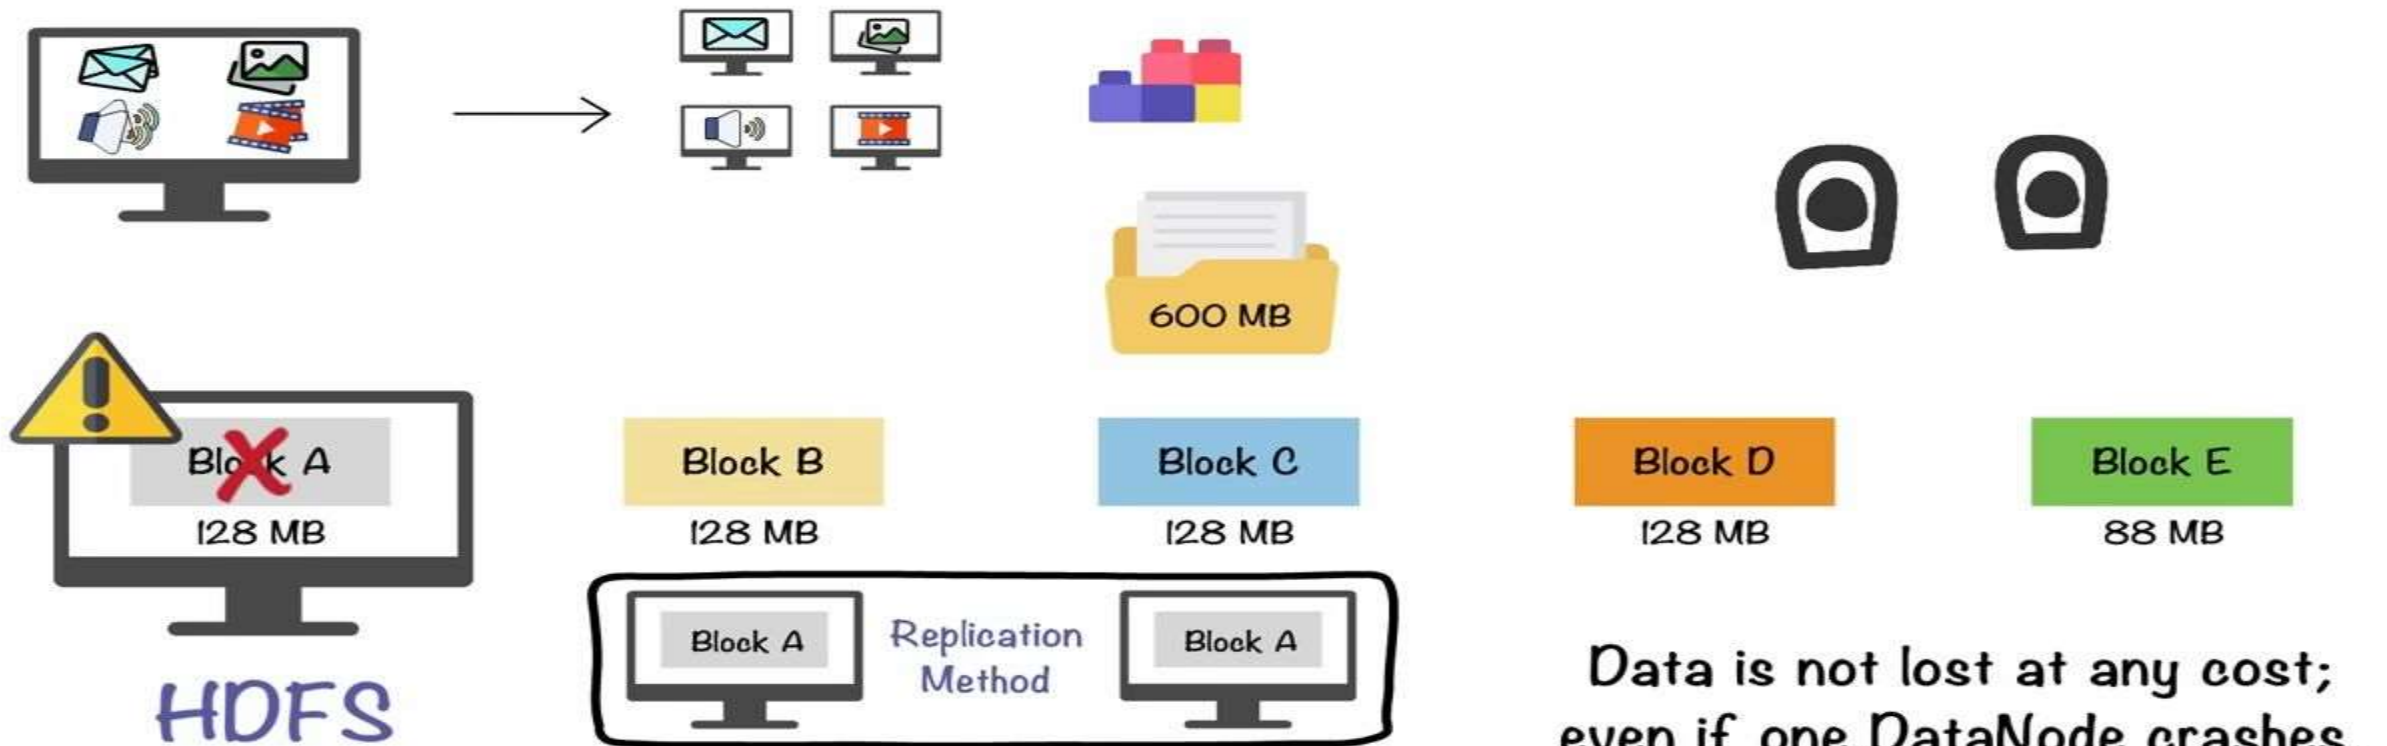




Hadoop consisted of three components that were specifically designed to work on big data

HDFS: Hadoop Distributed File System

I. Storage unit → HDFS




HDFS makes copies of the data and stores it across multiple systems

Data is not lost at any cost; even if one DataNode crashes, making HDFS fault-tolerant

MapReduce improves data processing efficiency by splitting large datasets into parts for parallel processing on multiple nodes, followed by aggregation of the results.

2. MapReduce

Traditional data processing method



Processing large volumes of a variety of data

Input

Split

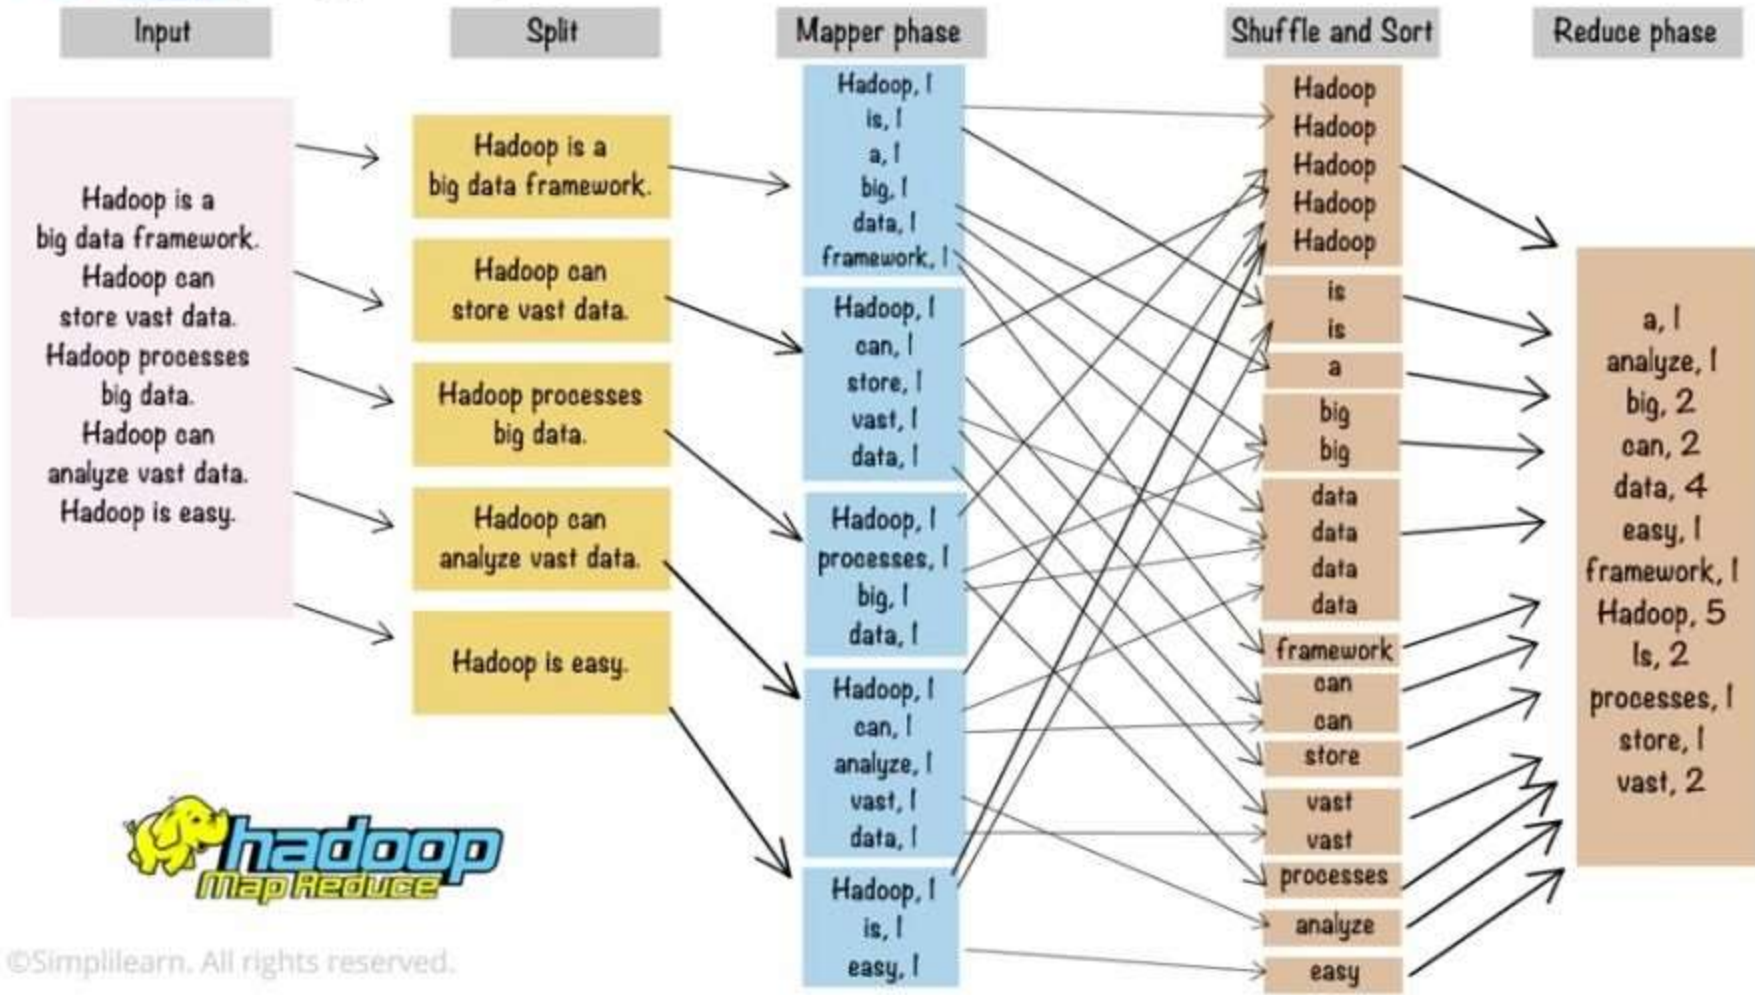


Mapper phase

Shuffle and Sort

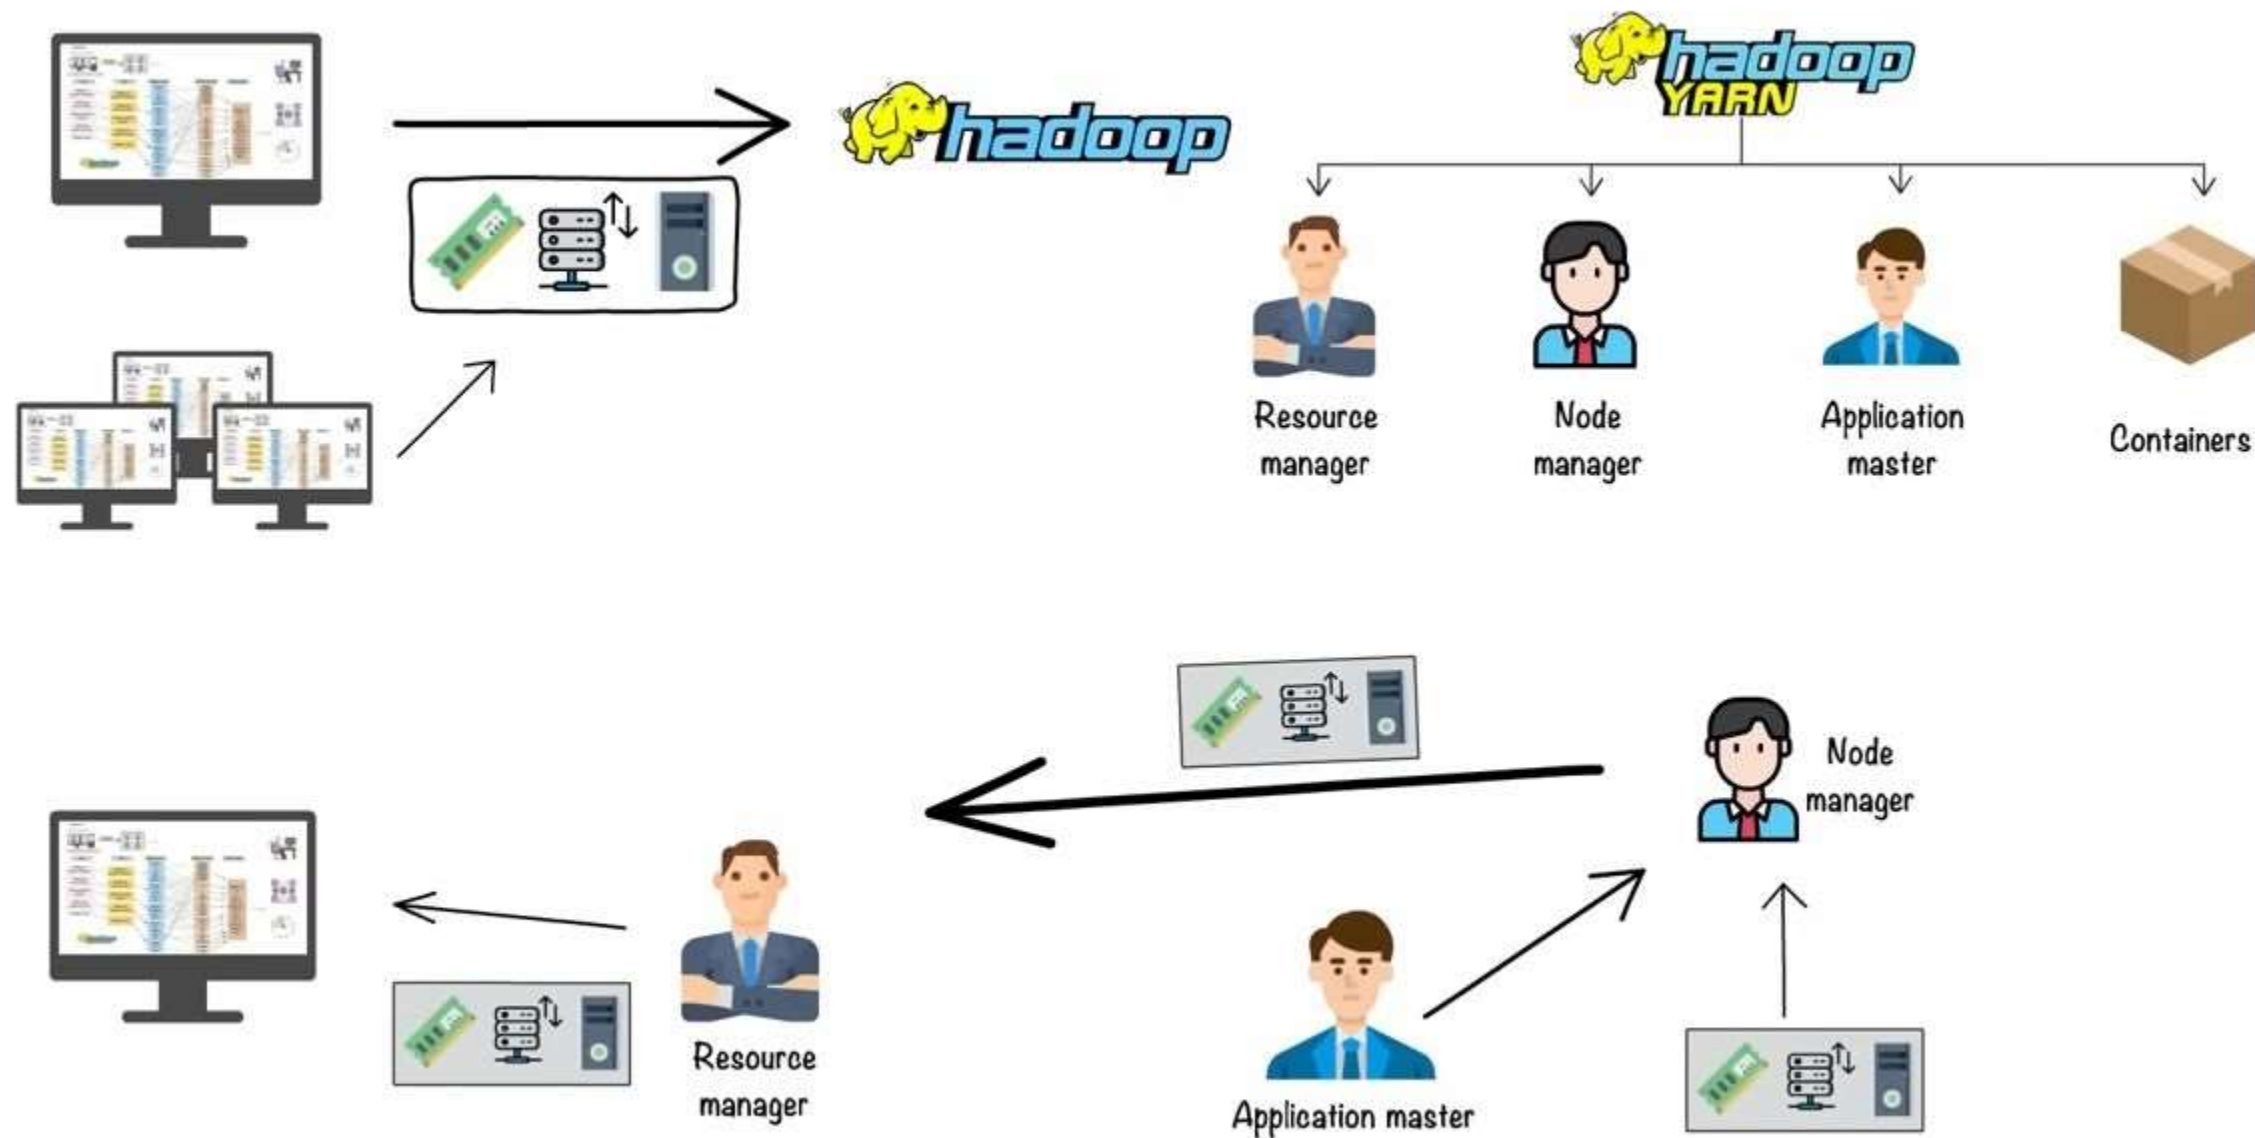
Reduce phase

2. MapReduce

Traditional data processing method



3. YARN



YARN processes job requests and manages cluster resources



What is the advantage of the 3x replication schema in HDFS?

- a) Supports parallel processing
- b) Faster data analysis
- c) Ensures fault tolerance
- d) Manages cluster resources